

18.065 Pset 1

Due Friday 2/17 at 1pm. Submit in PDF format: a decent-quality scan/image of any handwritten solutions (e.g. get a scanner app on your phone or use a tablet), and a PDF printout of your Jupyter notebook showing your code and (clearly labeled) results.

Problem 1 (4+4+4+4+4 points)

Recall from class that multiplying an $m \times p$ by a $p \times n$ matrix costs mnp scalar multiplications (and a similar number of additions) by the standard (practical) algorithms.

Matrix multiplication is **not commutative** ($AB \neq BA$ in general), but it **is associative**: $(AB)C = A(BC)$. It turns out that where you put the parentheses (i.e. in *what order* you do the multiplications) can make a *huge* difference in computational cost.

(a) If $x \in \mathbb{R}^n$ and A, B are $n \times n$ matrices, compare the scalar multiplication counts of $(AB)x$ vs. $A(Bx)$, i.e. if we do the multiplications in the order indicated by the parentheses.

(b) If $x, b \in \mathbb{R}^n$, **how many scalar multiplications** does the computation

$$p = (I - (xx^T)/(x^T x))b$$

take if we *do it in the order indicated by the parentheses*? (Note that dividing by a scalar α is equivalent to multiplying by α^{-1} at the negligible cost of one scalar division.)

(c) Explain how to compute the *same* p as in part (b) using as *few multiplications as possible*. Outline the sequence of computational steps, and give the count of multiplications.

(d) $p^T x =$ what?

(e) Implement your algorithm from (c) in Julia, filling in the code below, and time it for $n = 1000$ using the `@btime` macro from the [BenchmarkTools package](#), along with the algorithm from part (b), following the outline below. How does the ratio of the two times compare to your ratio of multiplication counts?

```
In [32]: using LinearAlgebra, BenchmarkTools

# algorithm from part (b)
function part_b(x, b)
    return (I - (x*x')*(x'*x)^-1) * b
end

# algorithm from part (c)
function part_c(x, b)
    return b - (x*((x'*(x'*x)^-1)*b))
end
```

```

# test and benchmark on random vectors:
n = 1000
x, b = rand(n), rand(n)

# test it first – should give same answer up to roundoff error
if part_c(x, b) ≈ part_b(x, b)
    println("Hooray, part (c) and part (b) agree!")
else
    error("You made a mistake: part (c) and part (b) do not agree!")
end

# benchmark it:

println("\npart (b): ")
@btime part_b($x, $b);

println("\npart (c): ")
@btime part_c($x, $b);

```

Hooray, part (c) and part (b) agree!

part (b):
5.651 ms (7 allocations: 22.90 MiB)

part (c):
2.044 μs (3 allocations: 23.81 KiB)

Problem 2 (8+4 points)

(a) Describe the four fundamental subspaces of the **rank-1 matrix** $A = uv^T$ where $u \in \mathbb{R}^m$ and $v \in \mathbb{R}^n$.

(b) For any column vectors $u, v \in \mathbb{R}^3$, the matrix uv^T is rank 1, except when _____, in which case uv^T has rank ____.

Problem 3 (5+4+4+4 points)

(a) Pick the choices that makes this statement correct for arbitrary matrices A and B : $C(AB)$ (contains / is contained in) the column space of (A / B) . Briefly justify your answer.

(b) Suppose that A is a 1000×1000 matrix of rank < 10 . Suppose we multiply it by 10 random vectors x_1, x_2, \dots, x_{10} , e.g. generated by `randn(1000)`. How could we use the results to get a 10×10 matrix C whose rank (almost certainly) matches A 's?

(c) Suppose we instead make 1000×10 matrix X whose columns are x_1, x_2, \dots, x_{10} . Give a formula for the *same* matrix C in terms of matrix products involving A and X .

(d) Fill in the code for C below, and compare the biggest 10 singular values of A (chosen to be rank ≈ 4 in this case) to the corresponding 10 singular values of C . Does it match what

you expect?

In [30]: `using LinearAlgebra`

```
# random 1000x1000 matrix of rank 4
A = randn(1000, 4) * randn(4, 1000)
@show svdvals(A)[1:10]

X = randn(1000, 10)
C = X' * A * X
@show svdvals(C)
```

```
(svdvals(A))[1:10] = [1033.437033249335, 1011.3839299777054, 967.3012768531058, 91
7.088555589723, 1.0221088602277808e-12, 8.667944633290469e-13, 7.802321942478929e-
13, 7.392352479625353e-13, 7.36152674855863e-13, 7.208168891076064e-13]
svdvals(C) = [18923.711378553286, 10128.014679660939, 5240.844393768108, 4083.3439
333669785, 3.5949835513878783e-12, 3.0672795174495276e-12, 2.5772864697684613e-12,
2.1910666302997e-12, 7.131643512506856e-13, 2.910189199817205e-13]
```

Out[30]: 10-element Vector{Float64}:

```
18923.711378553286
10128.014679660939
5240.844393768108
4083.3439333669785
 3.5949835513878783e-12
 3.0672795174495276e-12
 2.5772864697684613e-12
 2.1910666302997e-12
 7.131643512506856e-13
 2.910189199817205e-13
```

Problem 4 (4+5+5 points)

The famous Hadamard matrices are filled with ± 1 and have orthogonal columns (orthonormal if we divide H_n by $1/\sqrt{n}$). The first few are:

$$H_1 = (1), \quad (1)$$

$$H_2 = \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \quad (2)$$

$$H_4 = \begin{pmatrix} H_2 & H_2 \\ H_2 & -H_2 \end{pmatrix} = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & -1 \\ 1 & 1 & -1 & -1 \\ 1 & -1 & -1 & 1 \end{pmatrix}. \quad (3)$$

Notice that (for power-of-2 sizes), they are built up "recursively" out of smaller Hadamard matrices. Multiplying a vector by a Hadamard matrix requires no multiplications at all, only additions/subtractions.

(a) If you multiply $H_4 x$ for some $x \in \mathbb{R}^4$ by the normal "rows-times-columns" method (without exploiting any special patterns), exactly how many scalar additions/subtractions are required?

(b) Let's break x into two blocks: $x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$ for $x_1, x_2 \in \mathbb{R}^2$. Write out $H_4 x$ in terms of a sequence of 2×2 block multiplications with $\pm H_2$. You'll notice that some of these 2×2 multiplications are repeated. If we re-use these repeated multiplications rather than doing them twice, we can save a bunch of arithmetic — what is the new count of scalar additions/subtractions if you do this?

(c) Similarly, the 8×8 Hadamard matrix $H_8 = \begin{pmatrix} H_4 & H_4 \\ H_4 & -H_4 \end{pmatrix}$ is made out of H_4 matrices.

To multiply it by a vector $y \in \mathbb{R}^8$, the naive rows-times-columns method would require ____ scalar additions/subtractions, whereas if you broke them up first into blocks of 4, used your solution from (b), and then re-used any repeated H_4 products, it would only require ____ scalar additions/subtractions.

Problem 5 (5+5 points)

The famous "discrete Fourier transform" matrix F has columns that are actually eigenvectors of the (unitary) permutation matrix:

$$P = \begin{pmatrix} & 1 & & \\ & & 1 & \\ & & & 1 \\ 1 & & & \end{pmatrix}$$

for the 4×4 case, and similarly for larger matrices.

(a) One way of saying why Fourier transforms are practically important is that they *diagonalize* (are eigenvectors of) matrices that *commute* with P . If A is a 4×4 matrix whose first row is $(a \ b \ c \ d)$

$$A = \begin{pmatrix} a & b & c & d \\ ? & ? & ? & ? \\ ? & ? & ? & ? \\ ? & ? & ? & ? \end{pmatrix}$$

that commutes with P (i.e. $AP = PA$), what must be true of the other ("?") entries of A ?

(b) Fill in the matrix `A` in Julia below and fill in and run the code to check that it commutes with P and is diagonalized by F :

```
In [35]: a, b, c, d = 1, 7, 3, 2    # 4 arbitrarily chosen values

P = [0 1 0 0
      0 0 1 0
      0 0 0 1
      1 0 0 0]
```

```

F = im .^ ((0:3) .* (0:3)') # the 4x4 Fourier matrix

# fill in:
A = [a b c d
      d a b c
      c d a b
      b c d a]

```

```

Out[35]: 4x4 Matrix{Int64}:
 1  7  3  2
 2  1  7  3
 3  2  1  7
 7  3  2  1

```

```

In [36]: # check:
P * A == A * P

```

```

Out[36]: true

```

```

In [37]: # check that F diagonalizes A. (How?)

```

```

@show F'*A*F

```

```

F' * A * F = Complex{Int64}[52 + 0im 0 + 0im 0 + 0im 0 + 0im; 0 + 0im -8 + 20im 0
+ 0im 0 + 0im; 0 + 0im 0 + 0im -20 + 0im 0 + 0im; 0 + 0im 0 + 0im 0 + 0im -8 - 20i
m]

```

```

Out[37]: 4x4 Matrix{Complex{Int64}}:
52+0im  0+0im  0+0im  0+0im
 0+0im -8+20im  0+0im  0+0im
 0+0im  0+0im -20+0im  0+0im
 0+0im  0+0im  0+0im -8-20im

```